

Election Political Online Sentiment

PRESENTER: DENNIS HSU

CS 262

NOVEMBER 30, 2016

Outline

- Introduction
- Background
- Current Solution
- Randomization Techniques
- Implementation
- Experiment Results
- Conclusion/Future Work



Introduction

- Problem: Polls showing support for election political candidates are always changing and rarely accurate
- Polls are based on voluntary surveys
- Social Media, such as Twitter, provides an excellent source for data mining
 - Current Sentiment for each candidate
 - Not just voluntary
- Goal: To use Twitter to get the current online sentiment for political candidates (Clinton, Trump)

Introduction

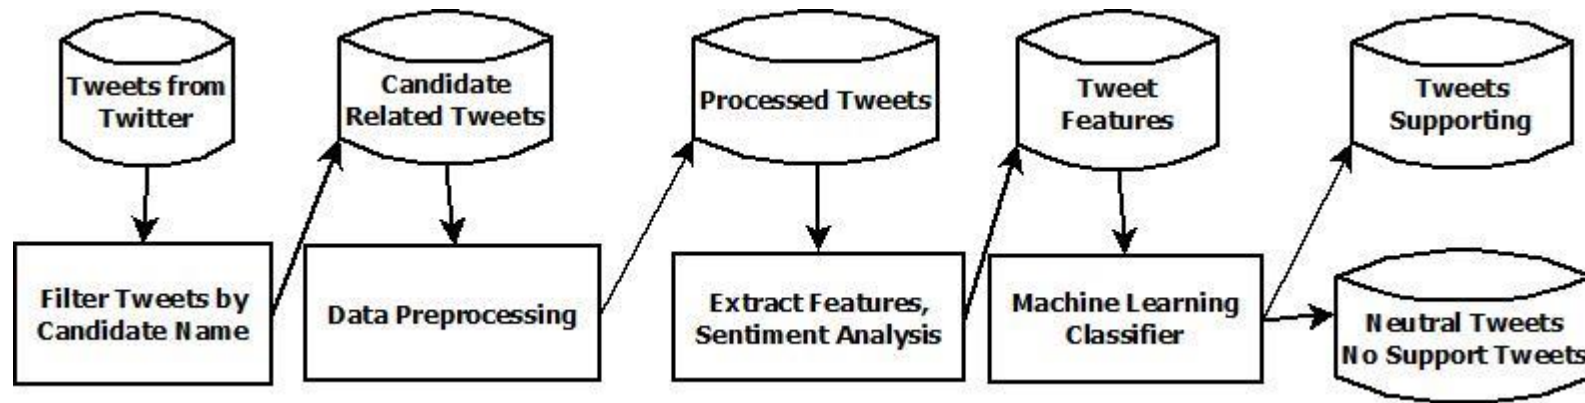
- Limitations
- Online Sentiment only
 - Vocal/Active Twitter users
 - “Echo Chamber” problem
- Online does not represent the silent majority
- Contribution: applied randomization to decrease amount of tweets stored to represent the current online sentiment/support for each political candidate

Background

- Techniques, Tools used
- Sentiment Analysis—used to label tweets as positive or negative
- Natural Language Processing—convert text/tweets into features
- Machine Learning—used to classify whether each tweet supports or not supports the candidate

Previous Solution

■ Pipeline(Before)



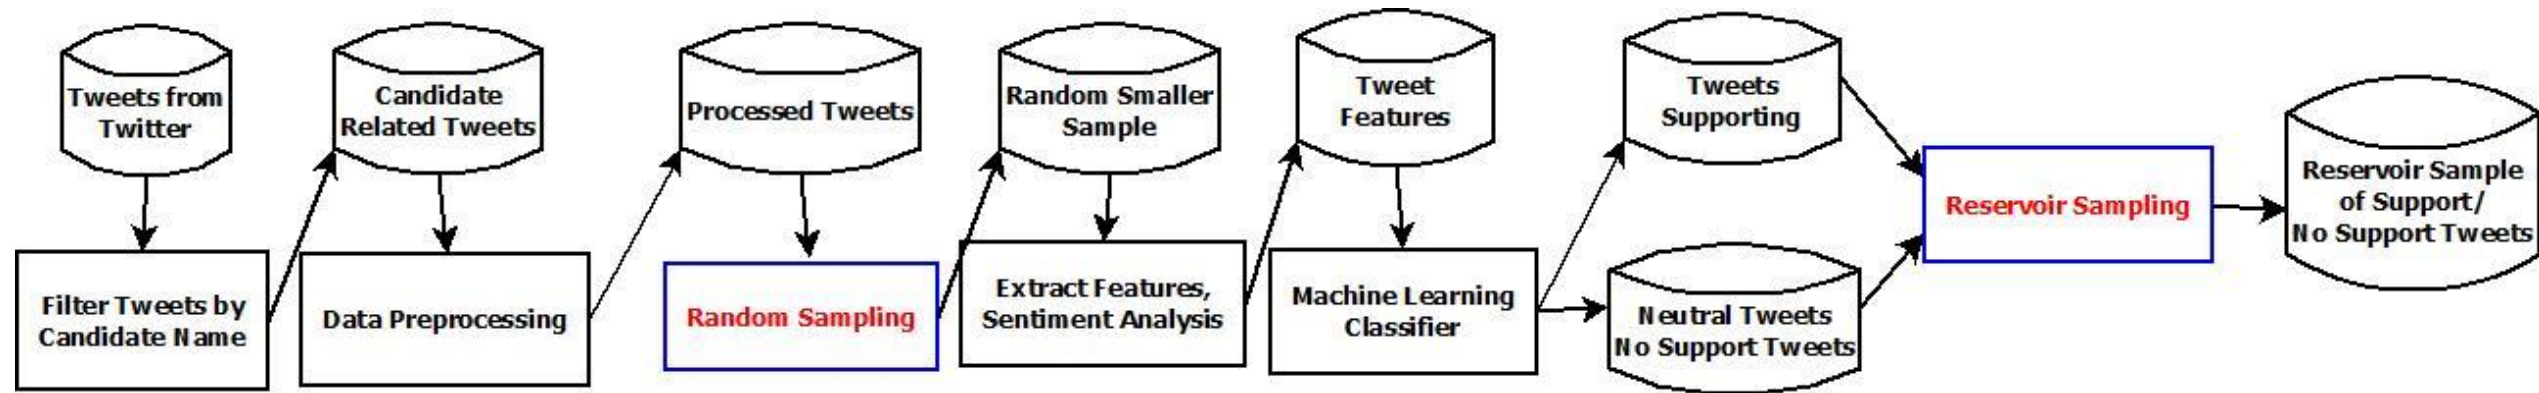
Current Solution

■Steps

1. Filter Tweets related to the candidate –(contains “Hillary Clinton” or “Donald Trump”)
2. Preprocess each Tweet into features—Sentiment Analysis, N-grams
3. Train Classifier separately for each candidate with labeled data (1 for support, 0 for oppose)
4. Use Train Classifier on an incoming stream of candidate-related tweets
5. Get current online sentiment for the candidate

Randomization Applied

■ Pipeline(Randomization)



Randomization Techniques

■ Randomization

1. Filter Tweets related to the candidate –(contains “Hillary Clinton” or “Donald Trump”)
2. Preprocess each Tweet into features—Sentiment Analysis, N-grams
3. Train Classifier separately for each candidate with labeled data (1 for support, 0 for oppose)
 - a) **Random smaller sample used instead of a larger set**
4. Use Train Classifier on an incoming stream of candidate-related tweets
5. Get current online sentiment for the candidate
 - a) **Reservoir Sampling—keep a smaller sample of tweets to represent larger set**

Randomization Techniques

- **Random smaller sample used instead of a larger set**
 - Decrease time to train classifier
 - Goal: Maintain accuracy of classifier
 - Expectation: Drop in accuracy of classifier
- **Reservoir Sampling—keep a smaller sample of tweets to represent larger set**
 - Online: Smaller Sample to keep to represent current set of classified tweets
 - Goal: Reduce amount of tweets needed to be kept
 - Expectation: Reservoir Sample should closely represent all tweets classified

Reservoir Sampling

- Keep the first tweet in memory
- When the i -th tweet arrives (for $i > 1$):
- with probability $1/i$, keep the new item(discard an old tweet)
- With probability $1 - 1/i$, keep old items(ignore new tweet)
- Induction:
 - when there is only one item, it is kept with probability 1
 - when there are 2 items, each of them is kept with probability $1/2$
 - when there are 3 items, the third item is kept with probability $1/3$, and each of the previous 2 items is also kept with probability $(1/2)(1-1/3) = (1/2)(2/3) = 1/3$
 - by induction, it is easy to prove that when there are n items, each item is kept with probability $1/n$
- From: https://en.wikipedia.org/wiki/Reservoir_sampling

```
(*  
  S has items to sample, R will contain the result  
*)  
ReservoirSample(S[1..n], R[1..k])  
  // fill the reservoir array  
  for i = 1 to k  
    R[i] := S[i]  
  
  // replace elements with gradually decreasing probability  
  for i = k+1 to n  
    j := random(1, i)  // important: inclusive range  
    if j <= k  
      R[j] := S[i]
```

Implementation

- My Laptop was stolen Dec.5th evening, lost data/results/powerpoint; had to redo it from scratch; All of the following slides were from the before results(will have update slides later)
- Tweepy—Filtered tweets related to “Hillary Clinton” and “Donald Trump”
 - 1,000,000 tweets per candidate filtered on election day (Nov 8, 2016) and day before (Nov 7)
- Dataset—“Hillary Clinton”
 - Labeled tweets for training
 - 402 tweets used to train classifier
 - Smaller random sample: 200 tweets used to train classifier
 - Balanced dataset: half positive, half negative/neutral

Implementation

- Tweet Pre-Processing to reduce noise, spam—Regex, Natural Language ToolKit(NLTK)
- Filter out tweets with links; filter out retweets
- Remove similar tweets
 - Similarity Index: 0.6 was found to be the best
 - Used Python's diff library SequenceMatcher
- Remove symbols, punctuation
- Tokenization into words
- Lemmatization—Base form of words to reduce noise

Implementation--Example

- Example filter
- Removed
 - RT @joshtpm: Hillary Clinton's popular vote lead now stands at 2.654 million votes, a 2 percentage point lead over Donald Trump, 48.2% to 4...
 - "CPAC 2013: Donald Trump: Immigration reform is a 'suicide mission' for GOP" <http://t.co/WdMLJcXZLL> by @SethMcLaughlin1
- Kept
 - 'A lot depends on who the real Donald Trump is.' - @BT_SDSC @PerthUSAsia #perthusasiatalks
 - a lot depends on who the real Donald Trump is perthusasiatalks

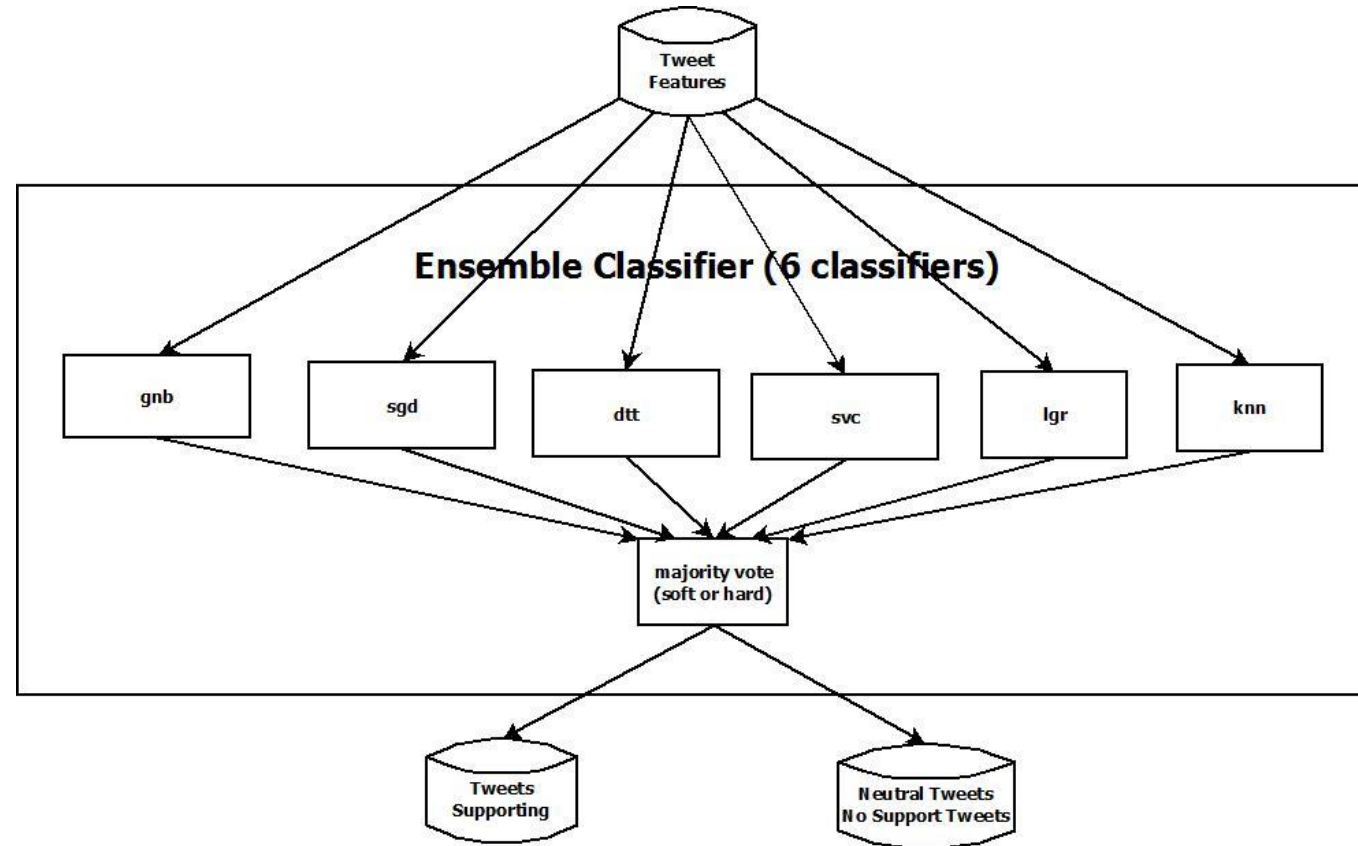
Implementation—Get Features

- Sentiment Analysis—lexicons—get a sentiment score for each tweet
 - Sentiwordnet—combine positive/negative scores
 - AFINN—get sum of scores
- N-grams—convert text into set of n-grams as features
 - Unigrams, bigrams—convert into sets of one word/two words as well as one character/two characters
- Combine together into set of features
- Convert features into a feature matrix using a vectorizer

Implementation--Classifier

- Train Classifier using feature matrix generated from labeled tweets
- Classifiers used
 - Gaussian Naïve Bayes(gnb)
 - Support Vector Classifier(svc)
 - Logistic Regression(lgr)
 - k-Nearest Neighbors(knn)
 - Decision Tree(dtc)
 - Stochastic Gradient Descent(sgd)
 - Random Forest Classifier(rfc)
 - Ensemble Classifier(gNB,SVC,LGR,kNN,DT,SGD)—(soft voc, hard voc)
 - Hard: majority vote
 - Soft: weighted probabilities of the vote
- Cross Validation(5-fold)—split dataset into 5 and trained/tested 5 times; collected f1 score

Ensemble Classifier



Implementation—Reservoir Sample

- Size of all tweets: 10,000
- Size of Reservoir Sample: 1,000
- Based on the algorithm in one of the previous slides, implementation was done
 - Loaded first 1000 tweets (if count < 1000)
 - Picked random number
 - If random number is less than or equal to 999, do replacement

```
r = random.randint(0, count)
```

```
if r < 999:
```

```
    #do prediction and replacement
```

- Compared with taking a random sample of 1000 from all tweets so far each time a new tweet comes in

Experiment Results—Old

- Stolen Laptop: Lost dataset, results, graphs, tweets of Hillary Clinton and Donald Trump from Election Days
- Data below was from what I remembered and presentation
- Testing of smaller sample used to classify: Hillary Clinton
 - 402 tweets: best classifier was Ensemble Classifier(~71-72%)
 - 200 tweets: best classifier was SVC(~71%)
- Testing of smaller sample kept: Hillary Clinton
 - From what I remember, Hillary had an approval rating of around 42% with full 10,000 tweets
 - Random Smaller Sample of 1000: approval rating varied above and below the full set, but error average was 1.8%
 - Reservoir Sampling of 1000: approval rating was below the full 10,000 tweets, but only by average error rate of 1.2%

Experiment Results—New(After Presentation)

- (Update: Since laptop was stolen, I have re-mined tweets again this time focusing on Donald Trump (Tweets from Dec. 6th) and his current support
- Labeled sample of 400 tweets (200 positive, 200 negative)
- Compare With smaller random sample of 200 tweets for training
- For reservoir sample testing: used 6000 tweets this time around of Donald Trump from Dec. 6, 2016
- Tested accuracy comparison of averages for 400 tweet-classifier on total, reservoir, and random sampling

Experiment Results—400 Sample Size; (f1 scores for each classifier)

| <u>count(word+score)(2-3) 93.146s</u> | <u>tfidf(char+score)(1-4) 102.172s</u> | <u>tfidf(word+score)(1-3) 107.248s</u> |
|---------------------------------------|--|--|
| svc: 0.649561568085 | svc: 0.689612222526 | svc: 0.684670519728 |
| gnb: 0.5208608176 | gnb: 0.601457339317 | gnb: 0.559498762542 |
| lgr: 0.641805527404 | lgr: 0.694846586696 | lgr: 0.679301918208 |
| sgd: 0.632294487482 | sgd: 0.579201178901 | sgd: 0.584666727607 |
| knn: 0.466596353505 | knn: 0.596530272476 | knn: 0.600343549055 |
| dtc: 0.571007872994 | dtc: 0.583951333026 | dtc: 0.528953643075 |
| rfc: 0.463463336887 | rfc: 0.575218866183 | rfc: 0.533389033991 |
| soft voc: 0.594174633452 | soft voc: 0.674683097813 | soft voc: 0.60144131876 |
| hard voc: 0.608975062333 | hard voc: 0.696358144974 | hard voc: 0.679381774203 |

Experiment Results—200 sample size

| <u>tfidf(word+score)(1-3) 15.493s</u> | <u>tfidf(char+score)(1-4) 21.516s</u> | <u>tfidf(char wb+score)(1-4) 13.687s</u> |
|---------------------------------------|---------------------------------------|--|
| svc: 0.519064566296 | svc: 0.569200893497 | svc: 0.555067332602 |
| gnb: 0.59594614386 | gnb: 0.635151969981 | gnb: 0.60772945497 |
| lgr: 0.520354460224 | lgr: 0.577686740766 | lgr: 0.540404076944 |
| sgd: 0.642713979644 | sgd: 0.536354228685 | sgd: 0.548920894017 |
| knn: 0.541998670405 | knn: 0.546332346768 | knn: 0.543829672625 |
| dtc: 0.472733477694 | dtc: 0.587154029445 | dtc: 0.566337112722 |
| rfc: 0.4963818842 | rfc: 0.491128224627 | rfc: 0.598614631165 |
| soft voc: 0.55935552889 | soft voc: 0.6133339599 | soft voc: 0.645151324792 |
| hard voc: 0.573636373611 | hard voc: 0.613755190686 | hard voc: 0.598973556753 |

Experiment Results—Reservoir

6000 Tweets(Total) Run time: 102.007 seconds

1000 Tweet Random Sampling runtime: 118.039 seconds

1000 Tweet Reservoir Sampling runtime: **54.332 seconds**

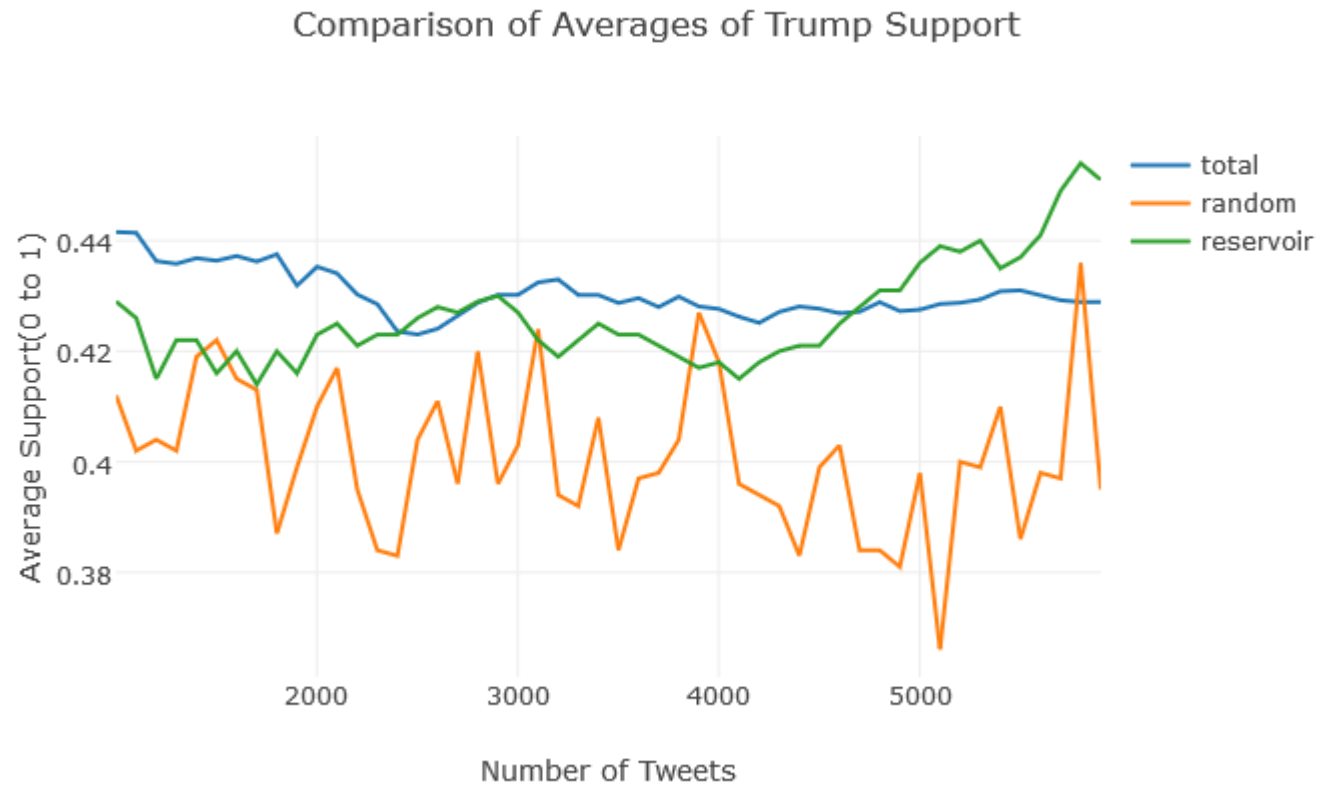
Experiment Results—Reservoir

Total versus Reservoir:

0.044179291646% avg. error

Total versus Random:

0.067820708354% avg. error



Conclusion & Future Work

- Have tested two random sampling techniques to improve time of pipeline
 - Used random smaller sample to train classifier to decrease time taken while maintaining accuracy—found out larger sample is better than smaller sample (obviously)
 - Used reservoir sampling to decrease amount of tweets needed to be kept (save space) as well as save time compared to taking a random sample each time
- Can be applied to other social media; pull dataset from Yelp, Facebook, etc.
- Can be used in other political elections or even on approval ratings for bills/propositions
- Future work
 - Apply to Donald Trump Tweets over more time
 - Apply to another elected official/candidate
 - Change classifier into 3 classes: positive, neutral, negative and retest pipeline and randomization

References

Fan Yu, Melody Moh, Teng-Sheng Moh. "Towards Extracting Drug-Effect Relation from Twitter: A Supervised Learning Approach." IEEE International Conference on Intelligent Data and Security to be held New York. April 2016. Accessed Nov 2016.

Liang Wu (MSCS), Teng-Sheng Moh and Natalia Khuri, "Twitter Opinion Mining for Adverse Drug Reactions," Proceedings of the 2015 IEEE International Conference on Big Data (BigData), Santa Clara, California, Oct. 2015, pp.1570-1574. Accessed Nov 2016.
<http://ieeexplore.ieee.org/xpl/articleDetails.jsp?arnumber=7363922>

Tweepy. Last Retrieved on Nov 2016. <http://www.tweepy.org>

"Reservoir Sampling". Wikipedia. Last Retrieved Dec 2016. https://en.wikipedia.org/wiki/Reservoir_sampling

Stefano Baccianella et. al. "SENTIWORDNET 3.0: An Enhanced Lexical Resource for Sentiment Analysis and Opinion Mining". Istituto di Scienza e Tecnologie dell'Informazione. Last Retrieved Nov 2016. <http://nmis.isti.cnr.it/sebastiani/Publications/LREC10.pdf> ,
<http://sentiwordnet.isti.cnr.it/>

Finn Årup Nielsen. "AFINN". Informatics and Mathematical Modelling, Technical University of Denmark. Last Retrieved Dec 2016.
http://www2.imm.dtu.dk/pubdb/views/publication_details.php?id=6010

W. B. Cavnar, and J. M. Trenkle, "N-Gram-Based Text Categorization", in Proceedings of SDAIR-94, 3rd Annual Symposium on Document Analysis and Information Retrieval, Las Vegas, NV, pp. 161-175, 1994. Accessed Jun 2016.
<http://odur.let.rug.nl/~vannoord/TextCat/textcat.pdf>

Code: <https://github.com/denniseh7/CS262>